

Qwen3.5-9B: truncate the network after layer L

pruned model = embeddings → layers 0...L → final norm → LM head; all later layers (mixer and MLP) removed

